

AI PROGRAM · CAPSTONE

Capstone Project

Cohort 6

Pick a project, ship an AI product in 4 weeks — at your own pace.

Individual by default · Group optional · **4-week build** · Self-paced phases

Build a real AI product — your way



Individual by default

Everyone owns their own project end to end. You can opt to team up in a small group if you prefer.



Phases, not dates

Progress through clear phases instead of fixed weekly deadlines. Move at the pace that suits you.



4 weeks to ship

You have 4 weeks total from kickoff to a working MVP. A rough guide is one phase per week.



Bar over speed

What matters is meeting the quality bar and clearing each phase — not racing to the end.

TASK: PICK YOUR CAPSTONE

Choose from 10 projects across 3 tiers

Pick one project that matches the effort you want to invest. Higher complexity means a bigger surface area to design, build and evaluate.



High complexity

4 projects ·


Multi-step systems: retrieval, agents, integrations.



Medium complexity

3 projects ·

Focused tools with one strong AI workflow.



Low complexity

3 projects ·

Tight, single-purpose AI features you can ship fast.

High complexity



Knowledge base search for PMs

A “Glean for PMs” that searches across docs, specs and tickets and answers questions with sources.



Support agent with RAG + triage

An assistant that answers from a company’s docs and routes or drafts replies to support tickets.



Autonomous research agent

Give it a topic; it searches, reads and returns a structured, cited report with key findings.



AI code-review assistant

Reviews GitHub pull requests and leaves inline comments on bugs, style and risky changes.

Medium complexity



Shopping assistant Chrome extension

A browser extension that compares products, summarises reviews and suggests the best buy on a page.



Meeting summariser

Turns a transcript into a summary, key decisions and action items with owners and deadlines.



Resume ↔ job-description matcher

Scores a resume against a JD, flags gaps and suggests tailored edits to improve the match.

Low complexity



Lead-qualification voice agent

A voice agent for a D2C brand that qualifies inbound leads with a short scripted conversation.



Flashcard & quiz generator

Upload notes or a PDF and get auto-generated flashcards and quiz questions to study from.



Review sentiment dashboard

Pulls product reviews for a D2C brand and visualises sentiment, themes and common complaints.

Solo, group, or bring your own



Work solo (default)

You own the full project — design, build, AI integration and evals. Best for going deep at your own pace and owning every decision.



Team up (optional)

Prefer to collaborate? Form a small group on the same project. Every member must contribute and clear the bar.



Bring your own project

Don't want one of the 10? You can propose your own, as long as it clears the quality bar:

- A side project you want to pursue
- An internal project from your company
- Same phases and bar still apply

THE GROUND RULES

What gets you disqualified



Missing phase progress

Stalling on a phase or not completing it within the 4-week window. Don't leave work to the last moment.



Falling short of the bar

Submitting work that doesn't pass the quality bar for the project.




If you're disqualified — it's not the end

You can take the capstone again in the next cohort. The capstone can be deferred at most twice — after that, future participation in capstone projects won't be allowed.

THE JOURNEY

Four phases from idea to MVP

Self-paced over 4 weeks — a rough guide is one phase per week.





PHASE 1 • SCOPE & RESEARCH

Scope & Research



Pick your project & format

Choose one project (or bring your own) and decide whether you're going solo or in a group.



Study the competitors

Research the top 3–5 players in the space. Note what they do well and where they fall short.



Define the problem & success metrics

Focus on ONE key problem. Define how you'll measure the value the product creates.

DELIVERABLE A clear problem statement, competitor notes, and a measurable definition of success.

WORKED EXAMPLE • SUCCESS METRICS

Meeting summariser

- Completeness — % of action items captured (target 95%+)
- Attribution accuracy — items assigned to the right person
- Conciseness — summary length vs transcript (<10%)
- Hallucination rate — info not in transcript (<2%)
- Latency — summary for a 30-min meeting (<15s)



PHASE 2 • DESIGN

Design the experience & data



End-to-end UI design flow

Map the full user journey through the product, screen by screen. This can change later.



Data model

Once the UI flow is set, design the data model — the entities and relationships behind it.

DELIVERABLE An end-to-end UI flow (review milestone) and a first-pass data model.



PHASE 3 · BUILD & AI INTEGRATION

Build & integrate AI



Build the product

Implement frontend, backend and any traditional AI you need to make the core flow work.



Integrate an LLM API

Wire OpenAI, Anthropic or Gemini into at least one use case. (If there's no AI use-case, skip this.)

DELIVERABLE A working build with at least one live LLM-powered use case.

WORKED EXAMPLE · LLM INTEGRATION

Meeting summariser

- Input: raw meeting transcript (text or audio-to-text)
- Prompt: “extract summary, key decisions, action items with owner & deadline, open questions”
- Output: structured JSON — summary, decisions, action_items[], open_questions[]



PHASE 4 · EVALUATE & ITERATE

Validate, ship MVP & evaluate



Surface & validate risks

List the risks in your product, then the fastest, cheapest way to validate each one.



Define & build the MVP

Scope the smallest product that tests your key risks (STAR format), then build it.



Set up evals & measure

Build an evals benchmark, run v1, and measure how far the metrics are from your targets.

DELIVERABLE An MVP, an evals benchmark, and a first set of metrics measured against the baseline.

WORKED EXAMPLE · TEST & BASELINE

Meeting summariser

- Run ~20 diverse sample transcripts through the product
- Manually verify each output against the transcript
- Compute the metrics to set your current baseline
- Compare v1 results to the targets and iterate

Terminal

<https://code.claude.com/docs/en/quickstart>

type 'Claude'

Download the PDF from Whatsapp (move it to your desktop)

take the path of PDF and run *cd path*

Create

Now go build something great.

Pick your project, move through the phases, and ship in four weeks.



Any doubts? Drop a message in the WhatsApp group.